Submit responses to all tasks which don't specify a file name to Canvas in a file called assignment7.{txt, docx, pdf, rtf, odt} (choose one of the formats). Submit all plots (if any) on Canvas. Do not zip your Canvas submission.

All *source files* should be submitted in the `HW07` subdirectory on the `main` branch of your homework git repo with no subdirectories. For this assignment, your `HW07` folder should contain `task1_cub.cu`, `task1_thrust.cu`, `task2.cu`, `count.cu` and `task3.cpp`.

Important note: All commands or code must work on *Euler* with the `nvidia/cuda/11.6.0.lua` module and `gcc/9.4.0` module loaded. Loading the modules is done via
`$ module load nvidia/cuda gcc/9.4.0`
This is because *Euler* may be currently experiencing an environment bug that requires you to use `gcc/9.4.0` for compiling `Thrust`- or `CUB`-related code. `gcc/9.4.0` may not be a requirement in other working environments.

We encourage you to test on *Euler* before you submit your homework.

Please submit clean code. Consider using a formatter like clang-format.

IMPORTANT: Before you begin, copy any provided files from `Assignments/HW07` directory of the ME759 Resource Repo. Do not change any of the provided files since these files will be overwritten with clean, reference copies when grading.

---

1. (40 pts) In HW05, you have implemented a reduction using the first add during load approach. In this task, you will compare the performance of `Thrust` and `CUB` with the previous GPU implementation by performing a scaling analysis for the reduction problem.

   a) Implement in a file called `task1_thrust.cu` the `Thrust` version of reduction. It's expected to do the following (some details about copying between host and device are not included here but should be implemented in your code when necessary):

      - Create and fill with random `float` numbers in the range [-1.0, 1.0] a `thrust::host_vector` of length `n`, where `n` is the first command line argument as below.
      - Use the built-in function in `Thrust` to copy the `thrust::host_vector` into a `thrust::device_vector`.
      - Call the `thrust::reduce` function to perform a reduction on the previously generated `thrust::device_vector`.
      - Print the result of reduction.
      - Print the time taken to run the `thrust::reduce` function in *milliseconds* using CUDA events.
      - Compile: `nvcc task1_thrust.cu -Xcompiler -O3 -Xcompiler -Wall -Xptxas -O3 -std c++17 -o task1_thrust`
      - Run by submitting a `sbatch` script (where `n` is a positive integer): `./task1_thrust n`
      - Example expected output:
        ```
        3141
        0.012
        ```

   b) Implement in a file called `task1_cub.cu` the `CUB` version of the reduction based on the code example from this link. Specifically, you should do the following.

      - Stick with the same device memory allocation pattern as the code example (`DeviceAllocate()` and `cudaMemcpy()`). Do not use unified memory.
      - Modify the example program so that the host array `h_in` has length `n` where `n` is the first command line argument as below, then fill in `h_in` with random `float` numbers in the range [-1.0, 1.0].

- Call the `DeviceReduce::Sum` function that outputs the reduction result to the output array.
- Print the reduction result
- Print the time taken to run the `DeviceReduce::Sum` function (the actual one, not the one that's used to find the size of temporary storage needed) in *milliseconds* using CUDA events. It's recommended to remove the debug function wrapper when gauging the performance of the reduction.
- Compile: `nvcc task1_cub.cu -Xcompiler -O3 -Xcompiler -Wall -Xptxas -O3 -std c++17 -o task1_cub`
- Run by submitting a `sbatch` script (where `n` is a positive integer): `./task1_cub n`
- Example expected output:
  `3141`
  `0.012`

c) On an *Euler compute node*:

- Run `task1_thrust` for value $n = 2^{10}, 2^{11}, \cdots, 2^{30}$ and generate a pattern of time vs. `n` in $\log - \log$ scale.
- Run `task1_cub` for value $n = 2^{10}, 2^{11}, \cdots, 2^{30}$ and generate a pattern of time vs. `n` in $\log - \log$ scale.
- Overlay the above two patterns on top of the plot you generated for HW05 `task1` in a file called `task1.pdf`. If you dropped HW05, then using information derived from this post or other Piazza posts is fine.

d) Comment on the performance of the three implementations you came up with above.

2. (35 pts)

    (a) Implement in a file called `count.cu` the function `count` as declared and described in `count.cuh`. Your `count` function should be able to take a `thrust::device_vector`, for instance, named `d_in` (filled by integers), and fill the output `values` array with the unique integers that appear in `d_in` in ascending order, as well as the output `counts` array with the corresponding occurrences of these integers. A brief example is shown below:

        • Example input: `d_in` = [3, 5, 1, 2, 3, 1]
        • Expected output: `values` = [1, 2, 3, 5]
        • Expected output: `counts` = [2, 1, 2, 1]

    Hints:

        • Since the length of `values` and `counts` may not be equal to the length of `d_in`, you may want to use `thrust::inner_product` to find the number of "jumps" (when `a[i-1] != a[i]`) as you step through the sorted array (the input array is not sorted, so you would have to do a sort using `Thrust` built-in function). You can refer to Lecture 18 and 19 for `thrust::sort` examples. There are other valid options as well, for instance, `thrust::unique`.
        • `thrust::reduce_by_key` could be helpful.

    (b) Write a test program `task2.cu` which does the following:

        • Create and fill with random `int` numbers in the range [0, 500] a `thrust::host_vector` of length `n` where `n` is the first command line argument as below.
        • Use the built-in function in `Thrust` to copy the `thrust::host_vector` into a `thrust::device_vector` as the input of your `count` function.
        • Allocate two other `thrust::device_vector`s, `values` and `counts`, then call your `count` function to fill these two arrays with the results of this counting operation.
        • Print the last element of `values` array.
        • Print the last element of `counts` array.
        • Print the time taken to run the `count` function in *milliseconds* using CUDA events.
        • Compile: `nvcc task2.cu count.cu -Xcompiler -O3 -Xcompiler -Wall -Xptxas -O3 -std c++17 -o task2`
        • Run by submitting a `sbatch` script (where `n` is a positive integer): `./task2 n`
        • Example expected output:
            `370`
            `23`
            `0.13`

    (c) On an Euler *compute node*, run `task2` for value $n = 2^5, 2^6, \cdots, 2^{24}$ and generate a plot of time vs. `n` in $\log - \log$ scale in a file called `task2.pdf`.

3. (25 pts) Write a C++ program in a file called `task3.cpp` which does the following:

- Launches four OpenMP threads.
- Prints out the number of threads launched, with the format `Number of threads:   x` (followed by a newline), where `x` is the total number of threads. This should be printed only once.
- Lets each thread introduce itself, with the print format `I am thread No.   i` (followed by a newline), where `i` is the thread number. Each thread should do that only once.
- Computes and prints out the factorial of integers from 1 to 8, `a!=b` (followed by a newline), where `a` is one of the 8 integers, and `b` is the result of `a!`. This should be done in parallel with all 4 threads.

How to go about it, and what the expected output looks like:

- Compile: `g++ task3.cpp -Wall -O3 -std=c++17 -o task3 -fopenmp`
- Run by submitting a `sbatch` script: `./task3`
- Example expected output (as you can see, the order matters not):
  ```
  Number of threads:  4
  I am thread No.  0
  I am thread No.  3
  I am thread No.  1
  I am thread No.  2
  3!=6
  5!=120
  4!=24
  6!=720
  7!=5040
  8!=40320
  1!=1
  2!=2
  ```

Important note: this problem is meant to get you started with OpenMP, it is not CUDA anymore, so the following changes need to be made to your `slurm` script:

- `#SBATCH --gres=gpu:1` should be removed since GPU is not required in this assignment.
- `#SBATCH --nodes=1 --cpus-per-task=4` (or `-N 1 -c 4` for short) should be added, which requests one node with 4 cores. In this course, `--cpus-per-task` should generally be no more than 20.