ME 759
High Performance Computing for Engineering Applications
Assignment 2
Due Thursday 09/22/2022 at 9:00 PM

Submit responses to all tasks which don't specify a file name to Canvas in a file called assignment2.{txt, docx, pdf, rtf, odt} (choose one of these formats). Submit all plots (if any) on Canvas. Do not zip your Canvas submission.

All *source files* should be submitted in the `HW02` subdirectory on the `main` branch of your `git` repo. Please use the name `HW02` exactly as shown here (both in terms of capitalization & name). Other names like `hw2`, `hw02`, `HW2` will not be recognized by the grading scripts. The `HW02` subdirectory should have no subdirectories. For this assignment, your `HW02` folder should contain `task1.cpp`, `scan.cpp`, `task2.cpp`, `convolution.cpp`, `task3.cpp`, and `matmul.cpp`.

All commands or code must work on *Euler* with no modules loaded unless specified otherwise (post a question on Piazza if the term *module* is confusing; or search the ME459 slides for the term). The commands may behave differently on your computer, so be sure to test on *Euler* before you submit.

Please submit clean code. Consider using a formatter like clang-format.

IMPORTANT: Before you begin, copy any provided files from `Assignments/HW02` directory of the ME759 repo.

---

1. a) Implement the `scan` function in a file called `scan.cpp` with signature defined as in `scan.h`. You should write an inclusive scan[1] on your own and not use any library scan functions.

   b) Write a file `task1.cpp` with a `main` function which (in this order)

      i) Creates an array of `n` random `float` numbers between -1.0 and 1.0. `n` should be read as the first command line argument as below.

      ii) Scans the array using your `scan` function.

      iii) Prints out the time taken by your `scan` function in *milliseconds*[2].

      iv) Prints the first element of the output scanned array.

      v) Prints the last element of the output scanned array.

      vi) Deallocates memory when necessary.

      - Compile command: `g++ scan.cpp task1.cpp -Wall -O3 -std=c++17 -o task1`
      - Run command (`n` is a positive integer; keep in mind: use Slurm, do not run on head node):
        `./task1 n`
      - Example expected output (followed by a newline):
        `0.06`
        `0.65`
        `87.3`

   c) On an *Euler* compute node (using a Slurm Job), run `task1` for each value $n = 2^{10}, 2^{11}, \cdots, 2^{30}$ and generate a plot `task1.pdf` (with axis labels) which plots the time taken by your algorithm as a function of `n`. This is called a scaling analysis.

   Feel free to post your scaling analysis plot in case you want to help other colleagues get an idea of what they should obtain at the end of this exercise.

---

[1] Given an array $[a_0, a_1, \cdots, a_{n-1}]$, the inclusive scan function produces the array $[a_0, a_0 + a_1, a_0 + a_1 + a_2, \cdots, a_0 + a_1 + \cdots + a_{n-1}]$. (In general, a scan can involve any other associative operation, not only +, like in this task.)

[2] Recall the document `timing.md`.

2. Convolutions[3] appear very prominently in image processing and in other fields like the numerical solution of partial differential equations, machine learning, etc.

a) Implement the `convolve` function in a file called `convolution.cpp` with signature as provided in `convolution.h`.

The operation that needs to be implemented is as follows:

$$g[x,y] = \sum_{i=0}^{m-1}\sum_{j=0}^{m-1} \omega[i,j]f[x+i-\frac{m-1}{2}, y+j-\frac{m-1}{2}] \ ,$$

where $x, y = 0, \cdots, n-1$; f is the original image, ω is the mask, and g is the result of the convolution; m is the dimension of the square matrix ω, where $m \geq 1$ is assumed to be an odd number. There are several ways of dealing with the boundaries, but here we will assume that $f[i,j] = 1$ if one of the following two conditions is not satisfied, and $f[i,j] = 0$ if neither of the following two conditions is satisfied:

$$\begin{cases} 0 \leq i < n \ , \\ 0 \leq j < n \ . \end{cases}$$

In other words, we will pad zeros for corners, and pad ones for edges (excluding the corners).

Example (uses $n = 4$ and $m = 3$):

$$f = \begin{bmatrix} 1 & 3 & 4 & 8 \\ 6 & 5 & 2 & 4 \\ 3 & 4 & 6 & 8 \\ 1 & 4 & 5 & 2 \end{bmatrix}$$

$$\omega = \begin{bmatrix} 0 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \end{bmatrix}$$

$$g = \begin{bmatrix} 3 & 10 & 10 & 10 \\ 10 & 12 & 14 & 11 \\ 9 & 7 & 14 & 14 \\ 5 & 11 & 14 & 4 \end{bmatrix}$$

b) Write a file `task2.cpp` with a `main` function which (in this order)

i) Creates an `n`×`n` `image` matrix (stored in 1D in row-major order) of random `float` numbers between -10.0 and 10.0. The value of `n` should be read as the first command line argument.

ii) Creates an `m`×`m` `mask` matrix (stored in 1D in row-major order) of random `float` numbers between -1.0 and 1.0. The value of `m` should be read as the second command line argument.

iii) Applies the `mask` to `image` using your `convolve` function.

iv) Prints out the time taken by your `convolve` function in *milliseconds*.

v) Prints the first element of the resulting convolved array.

vi) Prints the last element of the resulting convolved array.

vii) Deallocates memory when necessary via the `delete` function.

- Compile command: `g++ convolution.cpp task2.cpp -Wall -O3 -std=c++17 -o task2`
- Run command (where `n` and `m` are positive integers and `m` is an odd number; do not run on the head node):
  `./task2 n m`
- Example expected output (followed by a newline):
  `0.1`
  `1.5`
  `52.36`

---

[3]See here for more on convolutions in image processing, just note that we use a slightly different formulation.

3. Implement in a file called `matmul.cpp` the four functions with signatures and descriptions as in `matmul.h` to produce the matrix product $C = AB$. Pay attention to the argument types defined in `matmul.h`. For all of the cases, the array `C` that stores the matrix C should be reported in row-major order.

a) `mmul1` should have three `for` loops: the outer loop sweeps index `i` through the rows of `C`, the middle loop sweeps index `j` through the columns of `C`, and the innermost loop sweeps index `k` through; i.e., to carry out, the dot product of the i^{th} row `A` with the j^{th} column of `B`. Inside the innermost loop, you should have a single line of code which increments C_{ij}. Assume that `A` and `B` are 1D arrays storing the matrices in row-major order.

b) `mmul2` should also have three `for` loops, but the two innermost loops should be swapped relative to `mmul1` (such that, if your original iterators are from outer to inner `(i,j,k)`, then they now become `(i,k,j)`). That is the only difference between `mmul1` and `mmul2`.

c) `mmul3` should also have three `for` loops, but the outermost loop in `mmul1` should become the innermost loop in `mmul3`, and the other 2 loops do not change their relative positions (such that, if your original iterators are from outer to inner `(i,j,k)`, then they now become `(j,k,i)`). That is the only difference between `mmul1` and `mmul3`.

d) `mmul4` should have the `for` loops ordered as in `mmul1`, but this time around `A` and `B` are stored as `std::vector<double>`. That is the only difference between `mmul1` and `mmul4`.

e) Write a program `task3.cpp` that accomplishes the following:

- generates square matrices `A` and `B` of dimension at least 1000×1000 stored in row-major order.
- computes the matrix product $C = AB$ using each of your functions (note that you may have to prepare `A` and `B` in different data types so they comply with the function argument types). *Your result stored in matrix `C` should be the same no matter which function defined at a) through d) above you call.*
- prints the number of rows of your input matrices, and for each `mmul` function in ascending order, prints the amount of time taken in *milliseconds* and the last element of the resulting `C`. There should be nine values printed, one per line

  - Compile command: `g++ task3.cpp matmul.cpp -Wall -O3 -std=c++17 -o task3`
  - Run command: `./task3`
  - Sample expected output:
    ```
    1024
    1.23
    2.365
    1.23
    2.365
    1.23
    2.365
    1.23
    2.365
    ```

f) In a couple sentences, explain the difference that you see in the times for `mmul1`, `mmul2` and `mmul3` *when running on an Euler compute node*. What would explain the performance results you report? Be as specific as possible. Also comment on the difference or similarity you see between `mmul1` and `mmul4`.